

Introduction to Large Language Models

Marco Calamo

Laboratory of Advanced Programming



SAPIENZA
UNIVERSITÀ DI ROMA

Introducing AI, GenAI, LLMs

A quick&dirty introduction with the eyes of an IS engineer




SAPIENZA
UNIVERSITÀ DI ROMA

AI Timeline

Artificial intelligence (AI) is technology that enables computers and machines to simulate human intelligence and problem-solving capabilities

Source: <https://www.ibm.com/topics/artificial-intelligence>



Source: <https://www.adlittle.com/en/insights/report/generative-artificial-intelligence-toward-new-civilization>



SAPIENZA
UNIVERSITÀ DI ROMA


Symbolic and Subsymbolic AI

Symbolic AI


also called rule-based or traditional AI, it refers to an artificial intelligence approach that relies on a base of **knowledge represented in the form of symbols or rules**

Subsymbolic AI

it refers to the approach that focuses on creating intelligent behavior by simulating the behavior in biological systems. It deals with the processing and analysis of data using **numerical methods and statistical models** to solve problems without the need for explicit rules or algorithms



Symbolic and Subsymbolic AI



Source: <https://towardsdatascience.com/symbolic-vs-subsymbolic-ai-paradigms-for-ai-explainability-6e3982c6948a>



SAPIENZA
UNIVERSITÀ DI ROMA

History of Language Models

What is a language model?

A language model is a probabilistic model of a natural language^[1]. Since the 60s researchers in the field of Artificial Intelligence have tried to *teach* the natural language to the machine.


A few milestones:

- **Eliza:** is the world's first chatbot. It simulated conversation by using a pattern matching and substitution methodology that gave users an illusion of understanding on the part of the program (*Symbolic AI*)
- **RNN:** is a type of neural network characterized by direction of the flow of information between its layers (*Subsymbolic AI*)
- **LSTM:** is a type of RNN aimed at dealing with the vanishing gradient problem present in traditional RNNs^[2] (*Subsymbolic AI*)
- **Transformers:** the major breakthrough in the field that made possible large language models (*Subsymbolic AI*)

[1] [Jurafsky, Dan; Martin, James H. "N-gram Language Models". \(2024\)](#)

[2] [Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." *Neural computation* 9.8 \(1997\): 1735-1780.](#)

Source: [Evolution of Language Models: From Rules-Based Models to LLMs](#)



Large Language Models - Basics


Transformers

A **transformer** is a deep learning architecture developed by Google and based on the **multi-head self attention mechanism**^[1]. They have the advantage of having no recurrent units. It is originally composed by an **encoder** and a **decoder** part.

What makes a model “large”

The introduction of the transformer architecture made possible the training of models with **billions of parameters**. Bigger models as of today are reaching one trillion parameters.


Even though the transformer architecture has been improved over the years (mostly because of $O(n^2)$ computation complexity), **it is still the base of modern Large Language Models**.




Encoder and Decoder Self-Attention^[1]

[1] [Vaswani, Ashish, et al. "Attention is all you need." \(2017\)](#)

Visualizing Attention



Self-attention




Attention-matrix heatmap

Bahdanau, et al. 2015. Neural machine translation by jointly learning to align and translate. In Proc. ICLR.

Cross-attention

LLM - Architecture and Evolution



LLM Evolutionary Tree[1]



[1] [Yang, Jingfeng, et al. "Harnessing the power of llms in practice: A survey on chatgpt and beyond \(2024\)"](#)



LLM - Architecture and Evolution

Encoder-Only Architecture



The transformer encoder architecture takes in a sequence of tokens and produces a fixed-size vector representation of the entire sequence, which can then be used for **classification**



LLM - Architecture and Evolution

Encoder-Decoder Architecture



The transformer encoder-decoder architecture is used for tasks like **language translation**, where the model must take in a sentence in one language and output a sentence in another language



LLM - Architecture and Evolution

Decoder-Only Architecture

The transformer decoder architecture is used for tasks like **language generation**. The decoder takes in a fixed-size vector representation of the context and uses it to generate a sequence of words, with each word being conditioned on the previously generated words



Large Language Models - Pre-Training

Pre-Training or Training?

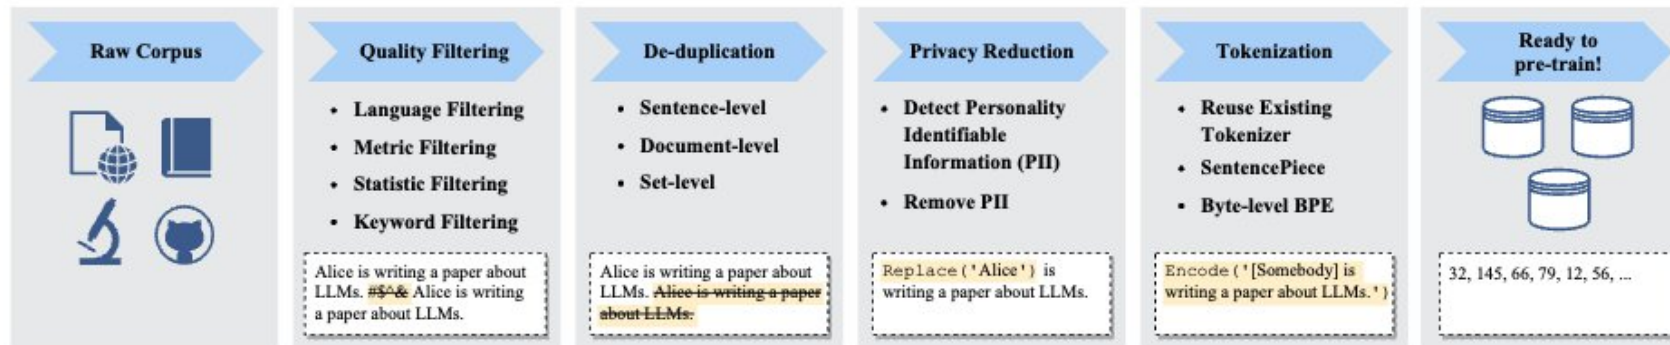
Pre-training is a foundational step in the LLM training process, where the model gains a general **understanding of language** by exposure to vast amounts of text data, in an unsupervised manner. It is called pre-training (instead of training) because after this stage the model is generally **not capable of solving specific problems**. Most pre-training common task are:

Masked Language Modeling (MLM): “The quick brown fox jumps over the ____ (lazy) dog.”

Next Sentence Prediction (NSP): “The scientist made a groundbreaking discovery. ____ (She was awarded the Nobel Prize).”

Contrastive Learning: “Hope is the thing with feathers — / That perches in the soul — / And sings the tune without the words — / And never stops at all — (Similar)”

Permutation Language Modeling (PLM): “the sky blue is so” -> “so blue is the sky”




[1] [Zhao, Wayne Xin, et al. "A survey of large language models." arXiv preprint arXiv:2303.18223 \(2023\)](#)

Pre-Training Pipeline[1]



Specializing Large Language Models



Optimizing a Language Model for a specific task

Source: <https://fractal.ai/blog/governing-llms-through-enhancing-techniques/>




SAPIENZA
UNIVERSITÀ DI ROMA

Fine-Tuning

After pre-training, LLMs can acquire the general abilities for solving various tasks. However, an increasing number of studies have shown that LLM's abilities can be further adapted according to specific goals with **fine tuning**.

Several optimization strategies can make the fine-tuning process **parameter efficient** and/or **memory efficient**



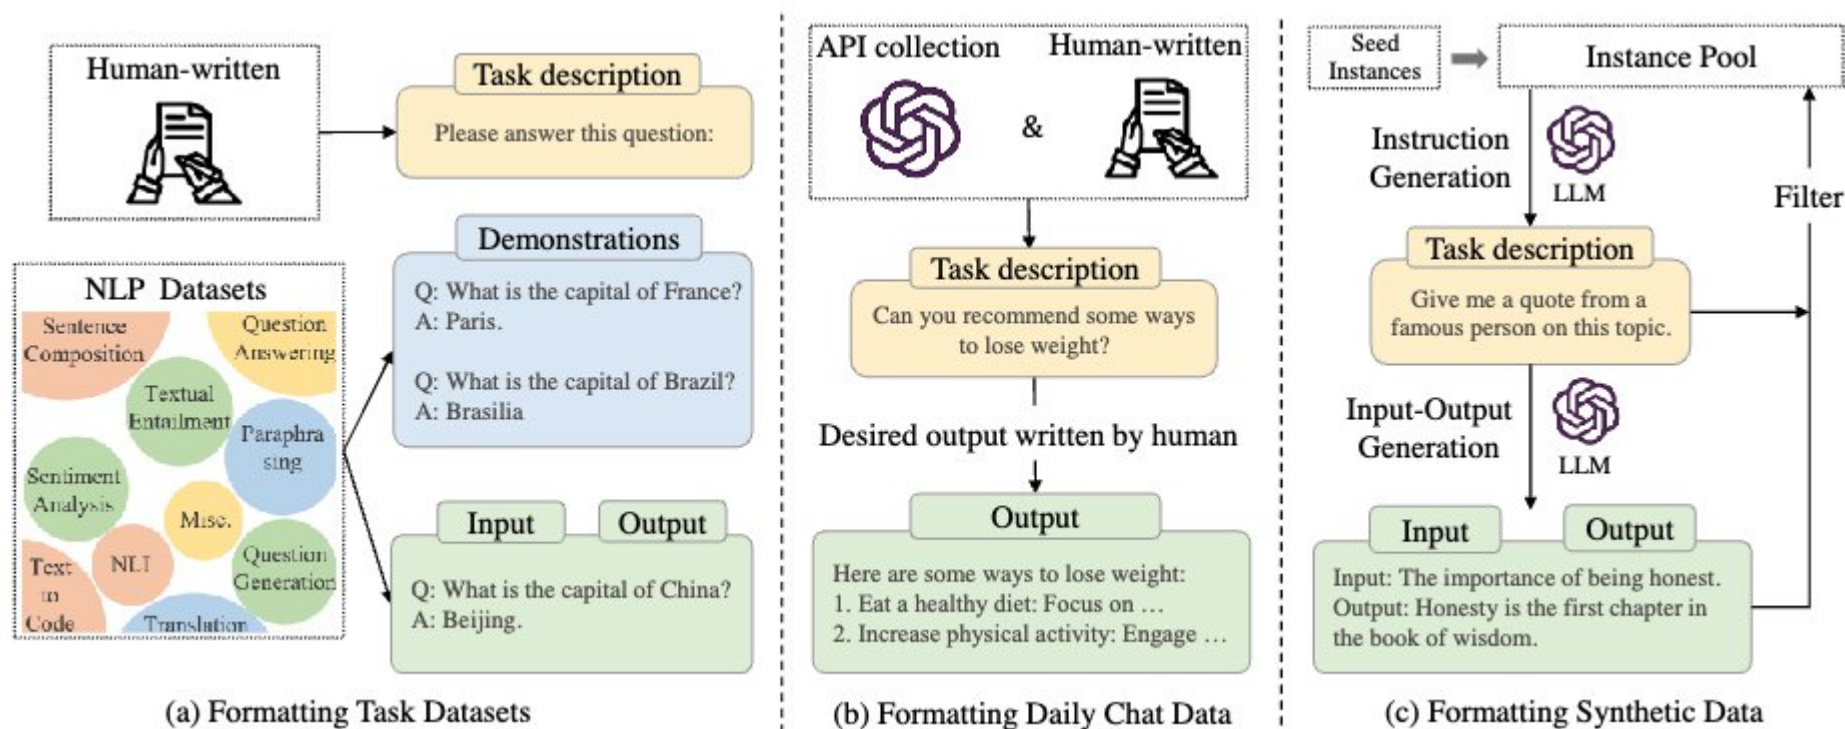
Source: <https://www.superannotate.com/blog/llm-fine-tuning>



SAPIENZA
UNIVERSITÀ DI ROMA

Fine Tuning - Instruction

Instruction tuning is the approach to fine-tuning pre-trained LLMs on a collection of formatted instances in the form of instructions in natural language. Even with a small dataset it can be very efficient



Instruction Fine tuning^[1]

[1] [Zhao, Wayne Xin, et al. "A survey of large language models." arXiv preprint arXiv:2303.18223 \(2023\)](#)



SAPIENZA
UNIVERSITÀ DI ROMA

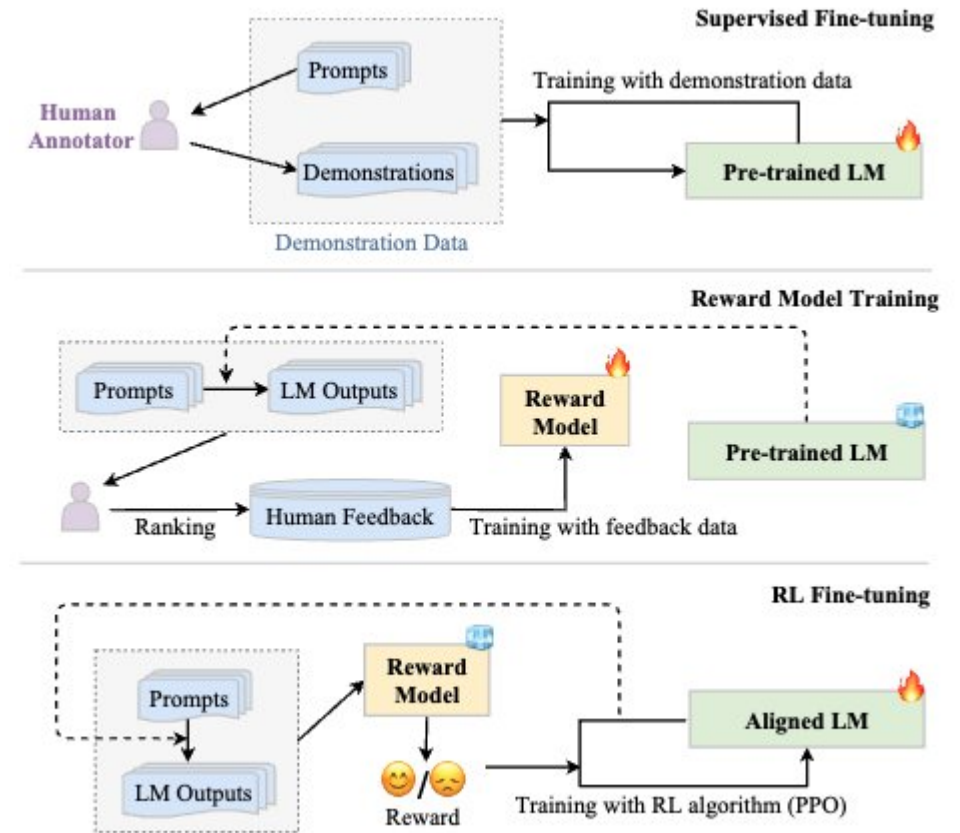
Fine Tuning - Alignment Tuning/RLHF

Alignment Tuning is used to avert these unexpected behaviors (e.g., fabricating false information, pursuing inaccurate objectives, and producing harmful, misleading, and biased expression)

Reinforcement Learning from Human Feedback (RLHF)

It is used to align LLMs with human values. It is a fine-tuned LLMs with the collected human feedback data, which is useful to improve the alignment criteria (e.g., helpfulness, honesty, and harmlessness). employing reinforcement learning (RL) algorithms to adapt LLMs to human feedback by learning a reward model.

- ChatGPT is an example



RLHF pipeline^[1]

[1] [Zhao, Wayne Xin, et al. "A survey of large language models." arXiv preprint arXiv:2303.18223 \(2023\)](#)

Fine Tuning limitations

Specialization is expensive

Fine-tuning is the **most effective** technique for specializing an LLM. However, it is **not always viable**. The major drawbacks of fine-tuning of LLM are **costs** in terms of:

- human resources (for both supervised and unsupervised training)
- time
- power
- CO2 emissions

Model		GPT-3	BLOOM	LLaMA	LLaMA-2	T5	PaLM
Developer		 OpenAI	 BigScience	 Meta		 Google	
Model Size (# parameters)		175B	175B	7B, 13B, 33B, 65B	7B, 13B, 34B, 70B	11B	540B
Training Data (# tokens)		300B	350B	1.4T	2T	34B	795B
Training Compute (FLOPs)		3.2E+23	3.7E+23	9.9E+23	1.5E+24	2.2E+21	2.6E+24
Processor	Manufacturer	Nvidia	Nvidia	Nvidia	Nvidia	Google	Google
	Type	GPU	GPU	GPU	GPU	TPU	TPU
	Model	V100	A100	A100	A100	TPU v3	TPU v4
Processor Hours		3,552,000	1,082,990	1,770,394	3,311,616	245,760	8,404,992
Grid Carbon Intensity (kgCO2e/KWh)		0.429	0.057	0.385	0.423	0.545	0.079
Data Center Efficiency (PUE)		1.1	1.2	1.1	1.1	1.12	1.08
Energy Consumption (MWh)		1,287	520	779	1,400	86	3,436
Carbon Emissions (tCO2e)		552	30	300	593	47	271

$KWh = \text{Hours to train} \times \text{Number of Processors} \times \text{Average Power per Processor} \times PUE \div 1000$
 $tCO2e = KWh \times \text{kg CO2e per KWh} \div 1000$

Source: <https://lajavaness.medium.com/llm-large-language-model-cost-analysis-d5022bb43e9e>



Fine Tuning - Parameter Efficient


Parameter-efficient fine-tuning has been an important topic that aims to reduce the number of trainable parameters while retaining a good performance as possible. Several techniques can be used:

Adapter Tuning: incorporates **small neural network modules** (called adapter) into the Transformer models. Only adapter parameters are updated during fine-tuning

Prefix Tuning: prepends a sequence of prefixes, which are a set of **trainable continuous vectors**, to each Transformer layer in language models. Only prefix parameters are updated during fine-tuning

Prompt Tuning: incorporates **trainable prompt vectors** at the input layer

Low-Rank Adaptation (LoRA): imposes the **low-rank constraint** for approximating the update matrix at each dense layer, to reduce the trainable parameters for adapting to downstream tasks



Parameter efficient fine tuning[1]

[1] [Zhao, Wayne Xin, et al. "A survey of large language models." arXiv preprint arXiv:2303.18223 \(2023\)](#)



SAPIENZA
UNIVERSITÀ DI ROMA


Fine Tuning - Memory Efficient

The most used techniques to reduce a model's memory footprint are either reducing floating points weight resolution from the full 32-bit down to 16-bit or **quantization**


Quantization

It maps weights from floating-point numbers to integers, of different resolutions, from 8-bit down to 2-bit


32-bit float (FP32)



16-bit float (FP16)




bfloat16 (BF16)



Source:

<https://towardsdatascience.com/introduction-to-weight-quantization-2494701b9c0c>



Source: <https://levelup.gitconnected.com/cant-afford-the-expensive-gpu-for-your-ai-development-try-4-bit-quantization-b26108229b75>



Prompt Engineering

Prompt engineering is the process of structuring an instruction that can be interpreted and understood by a generative AI model^[1]. It enables:

Zero-Shot


The prompt used to interact with the model **won't contain examples** or demonstrations

In-Context Learning or Few-Shot Learning

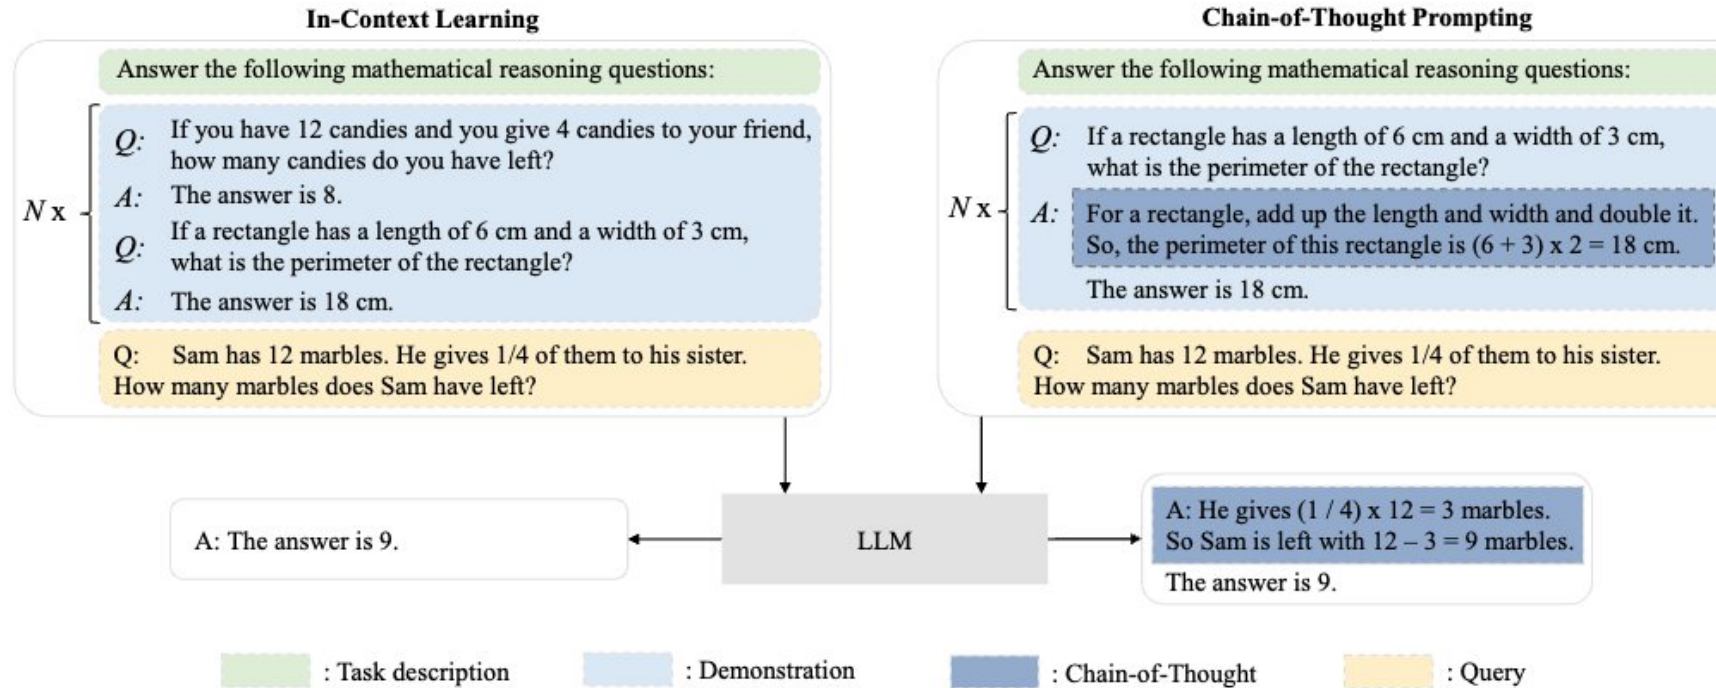
Demonstrations of the task are provided to the model as part of the prompt

Chain of thought/reasoning

Incorporates intermediate **reasoning steps**, which serve as the bridge between inputs and outputs



Prompt Engineering at Work




Source: [Zhao, Wayne Xin, et al. "A survey of large language models." arXiv preprint arXiv:2303.18223 \(2023\)](#)



Retrieval Augmented Generation (RAG)

RAG takes an input and retrieves a set of **relevant/supporting documents** given a source. The documents are **concatenated as context** with the original input prompt and fed to the text generator which produces the final output.




Source:

<https://medium.com/@bijit211987/advanced-rag-for-llms-slms-5bcc6fbba411>



RAG at Work



Tested with [meta-llama/Meta-Llama-3-8B-Instruct](#)



SAPIENZA
UNIVERSITÀ DI ROMA

RAG at Work

Home » Users » Massimo Mecella

MASSIMO MECELLA



Professore ordinario
stanza: B211
telefono: +39 0677274028

Biography

Massimo Mecella (born in Rome, Italy in 1973) holds a PhD in Computing Science and Engineering from SAPIENZA (earned on 21 January 2002). Currently he is associate professor in the School of Information Engineering, Computer Science and Statistics of SAPIENZA, in the Department of Computer, Control and Management Sciences & Engineering ANTONIO RUBERTI (DIAG, formerly DIS - Department of Computer and Systems Sciences). He has recently achieved (April 2017) the Italian National Qualification (ASN) to Full Professor. From March 2005 to August 2006 he has been a Postdoc on *Service Oriented Computing: Synthesis and Orchestration of Composite Services* at DIS. In the winter semester 2005, he has been a temporary research assistant at CERIAS (Center for Education and Research in Information Assurance and Security) and Department of Computer Science, Purdue University, USA, hosted by Elisa Bertino. From November 2002 to October 2004 he has been a research fellow in SAPIENZA (in the Latina campus). During the PhD period, in 1999 he has been *internship student* at Telcordia Applied Research, Telcordia Technologies (formerly Bellcore), Morristown, New Jersey, USA, hosted by Amjad Umar, Paolo Missier and Francesco Caruso.

Retrieval from the internet

Much better 🙏

Massimo Mecella is a **Professor of computer science** at the Sapienza University of Rome, Italy, and leader of the Processes, Services and Software Management group.

Who is Massimo Mecella?

User Prompt

LLM + RAG

Final Output

You are an assistant for question-answering tasks. Use the following pieces of retrieved context to answer the question. If you don't know the answer, just say that you don't know. Use three sentences maximum and keep the answer concise.

Question: Who is massimo mecella?


Context: Massimo Mecella | Department of Computer, Control and Management Engineering [...]

Rag generated prompt

Tested with [meta-llama/Meta-Llama-3-8B-Instruct](#) and [langchain v0.2](#)

Prompt Chaining

To improve the reliability and performance of LLMs, one of the important prompt engineering techniques is **to break tasks into its subtasks**. Once those subtasks have been identified, the LLM is prompted with a subtask and then its response is used as input to another prompt.



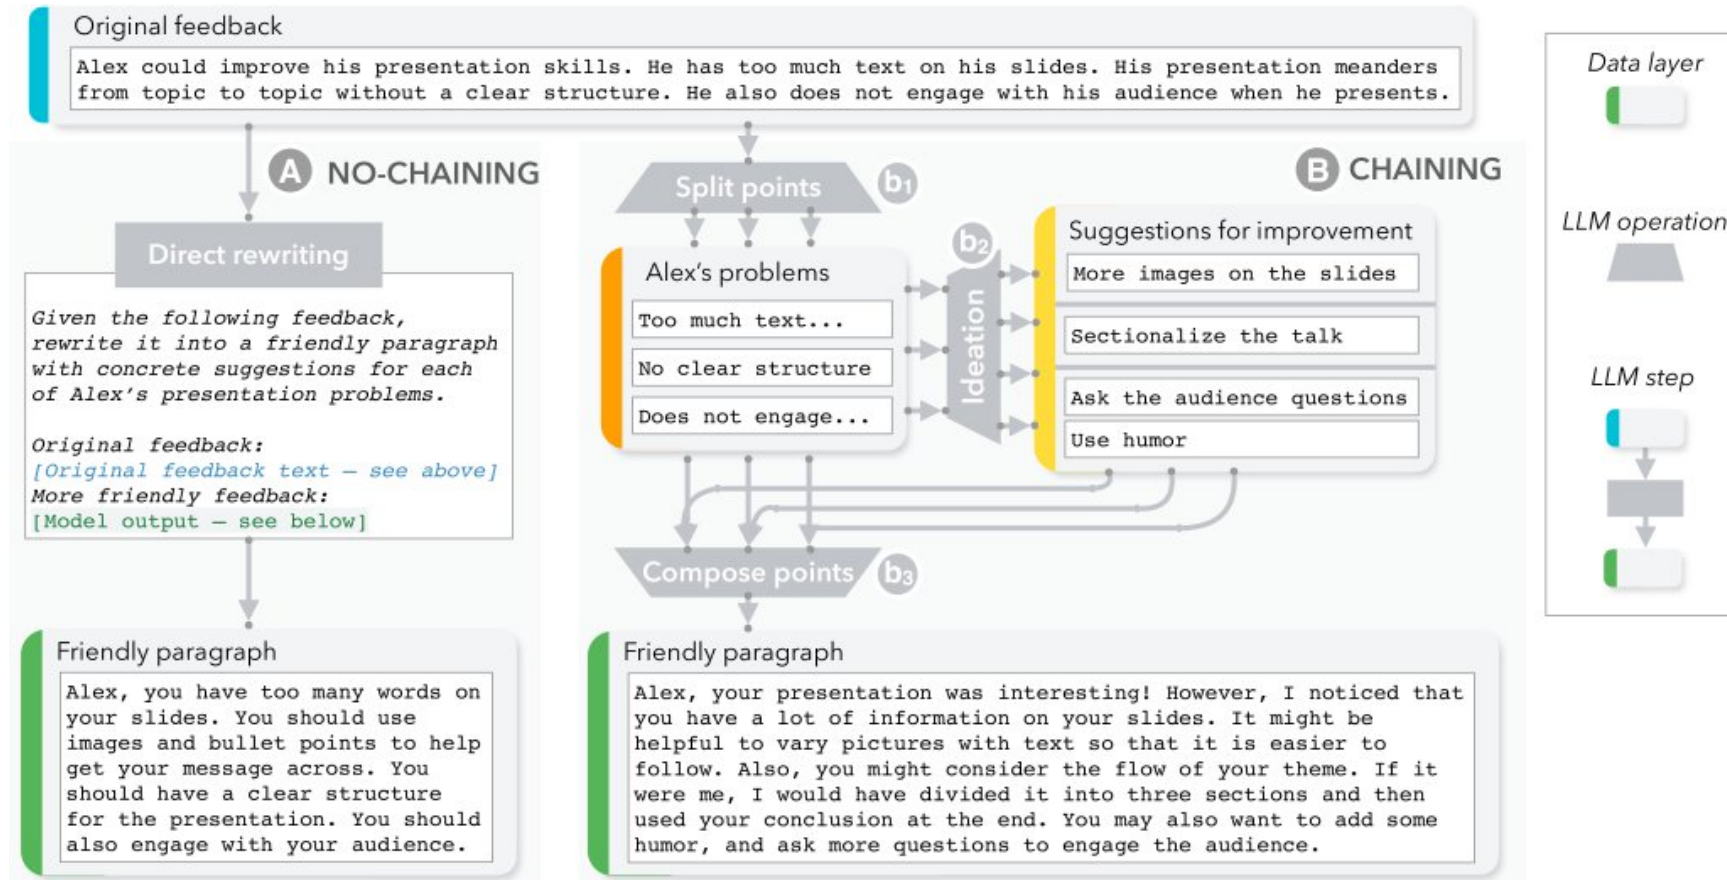
Source:

<https://aigenealogyinsights.com/2023/06/13/the-power-of-a-i-genealogical-prompt-chaining/>



SAPIENZA
UNIVERSITÀ DI ROMA

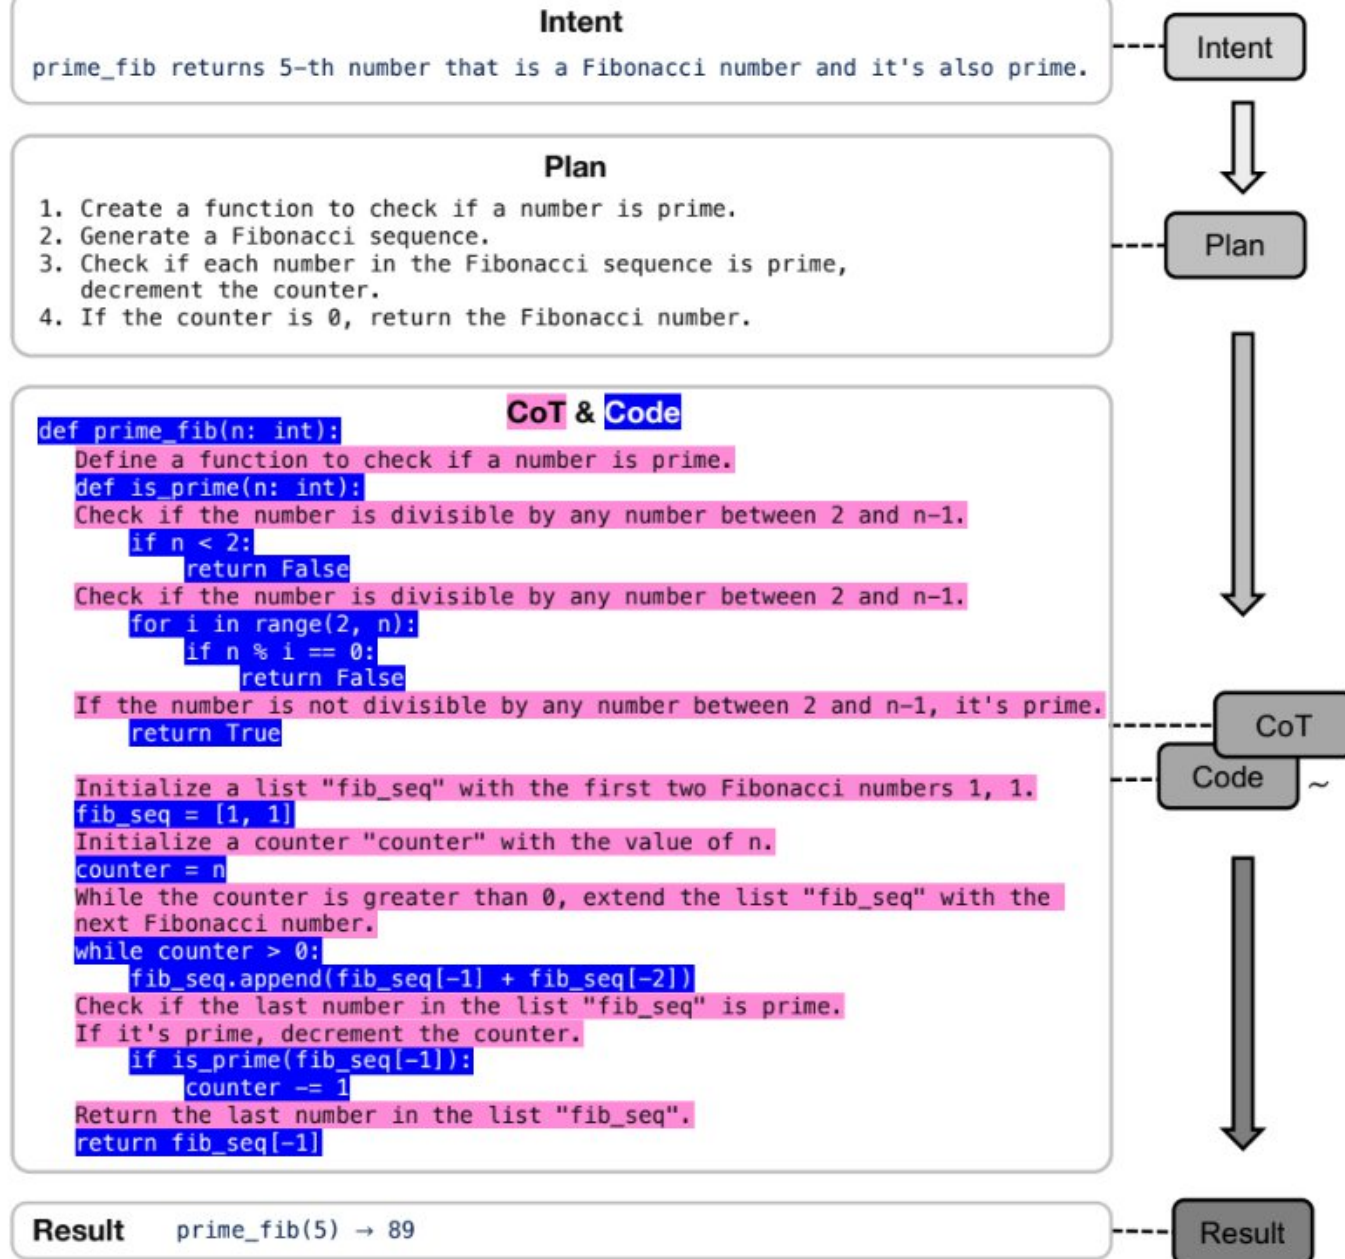
Prompt Chaining at Work



Source:


<https://cobusgreyling.medium.com/chaining-large-language-model-prompts-81091daccaad>






LLM Agent

LLM agents involve LLM applications that can execute **complex tasks** using an architecture that combines LLMs with key modules like **planning and memory**. When building LLM agents, an LLM serves as the **main controller** or "brain" that controls a flow of operations needed to complete a task or user request



Source: <https://www.promptingguide.ai/research/llm-agents>


```
agent.run("Caption the following image", image=image)
```



Source: https://huggingface.co/docs/transformers/transformers_agents




LLM Agent at Work



Source: <https://www.promptingguide.ai/research/llm-agents.en>



LLM Agent at Work II



AI Agents for Code Development

Github Copilot

GitHub Copilot is an AI-powered pair programmer that acts as a coding assistant to help developers write code faster and more efficiently. It provides context-aware code suggestions as you type, can write code from natural language prompts, and assists with various tasks like debugging, creating unit tests, and explaining code



Windsurf

Windsurf AI is an AI-powered integrated development environment (IDE) and coding assistant that enhances developer productivity through features like code generation, intelligent suggestions, and multi-file editing



Cursor

Cursor AI is a code editor that integrates AI directly into a familiar VS Code-like interface to help developers write code faster using natural language

