

X86 ASSEMBLY

MALWARE ANALYSIS AND INCIDENT FORENSICS
M.Sc. in Cyber Security

MALWARE ANALYSIS
M.Sc. in Engineering in Computer Science

A.Y. 2026/2027



SAPIENZA
UNIVERSITÀ DI ROMA



CIS SAPIENZA
CYBER INTELLIGENCE AND INFORMATION SECURITY

ANALYZING BINARY CODE

Daily routine for malware analysts 🤔

```
1 FDX 12:01a 23- 1
A 002000 C2 30 REP #$30
A 002002 18 CLC
A 002003 F8 SED
A 002004 A9 34 12 LDA #$1234
A 002007 69 21 43 ADC #$4321
A 00200A 8F 03 7F 01 STA $017F03
A 00200E D8 CLD
A 00200F E2 30 SEP #$30
A 002011 00 BRK
A 2012

r
PB PC NUmxDIZC .A .X .Y SP DP DB
; 00 E012 00110000 0000 0000 0002 CFFF 0000 00
g 2000

BREAK

PB PC NUmxDIZC .A .X .Y SP DP DB
; 00 2013 00110000 5555 0000 0002 CFFF 0000 00
m 7f03 7f03
>007F03 55 55 00 00 00 00 00 00 00 00 00 00 00 00 00 00:UU.....
```

ANALYZING BINARY CODE

Reverse engineering of binary code is difficult

Thankfully, only some portions of a malware sample really require it



FACT: malicious code $\neq$ compiler-generated code

- obfuscation of control and data flows
- subtle semantics (e.g., an instruction jumps into the middle of another)
- hand-written sequences (e.g., API hashing algorithms)
- self-modifying code
- other decoys

DECOMPIlation?

- Decompilers (e.g., Ghidra, Hex-Rays) try to reconstruct a high-level C-like representation of a binary code fragment
- Process far from perfect
 - information is lost during compilation (e.g., types, identifiers)
 - decompiling obfuscated code returns an obfuscated source
 - anti-analysis techniques
- But very helpful for reverse-engineering professionals

CPU AND MEMORY BASICS




ISA AND X86

- **Instruction Set Architecture** (ISA) = abstract model describing what a programmer should know to program the machine
- ISA acts as an interface between software and hardware
 - Humans or compilers produce ISA-conformant **assembly** code, then an assembler encodes it in a **binary** format understood by hardware
- Intel defined the **x86** ISA for its 8086 processor (1976-1978)
 - Later extensions preserved backward compatibility
 - Today most malware is still 32-bit
 - x86 is also known as IA-32
 - x86-64 (x64 for short) is the 64-bit successor

WORD SIZE AND ENDIANNESS

- A **word** is the natural data unit for a specific CPU design. For x86 processors, the word size is 32 bits
 - However, for backward compatibility, x86 instructions assume a **word** operand to be 16-bit long, while a **dword** operand is 32-bit long
 - IA32 common data types: byte, word, dword
- Memory always operates at byte level. **Endianness** specifies how multi-byte sequences are read from/written to memory




Say we want to store **dword** `0xABADCAFE` at address `0x1000...`

- **Big Endian**: most significant byte first
 - **Little Endian**: least significant byte first
- x86 uses Little Endian encoding

REGISTERS

- Coders use (many) variables. CPUs work on (few) **registers**
- Registers hold data with fast access time
 - General-purpose registers (GPRs) can store data or memory addresses
 - Status register holds truth values on the state of the processor after the last executed instruction (used, for example, in conditional computations)
 - Program counter holds the address of the instruction being executed
 - Many other registers
- x86 general-purpose registers are 32-bit wide

ACCESSING REGISTERS



Credits: Yale's CS421

8 general-purpose registers

- A, B, C, D, SI, DI
- ESP points to top of the stack
- EBP often set $:=$ ESP when a function begins (otherwise is available)

Prefix **E**AX indicates 32-bit size. When E is absent, only the least significant 16 bits (2 bytes) get manipulated. The two least significant bytes of A, B, C, and D can also be accessed individually

Program counter EIP and status register EFLAGS cannot be accessed directly

INSTRUCTION CYCLE


- Fetch-decode-execute cycle

- Fetch: CPU reads instruction from address stored in EIP
- Decode: control unit determines meaning of the instruction
- Execute: carry out computation using ALU or move data
- Then EIP advances to the next adjacent instruction, unless current instruction explicitly altered the control flow (e.g., with a jump)

- Instructions vary in size (1-15 bytes) and are stored consecutively in memory along with any **immediate operands** (i.e., data or addresses) they might use. Register operands are embedded instead in the **opcode** binary representation:

b8	01	00	00	00	mov	eax,	0x1
bf	01	00	00	00	mov	edi,	0x1

MEMORY MAP OF A PROCESS



Simplified view of a 32-bit Windows program

- lower 2GB for user space, upper 2GB (not shown) for kernel code
- program image = PE header + .text + .rdata + .data + .rsrc
- **stack grows to low addresses**
- DLLs have their own areas
- fixed addresses for first Thread Entry Block, Process Entry Block, and a special page that the kernel shares with userland

MEMORY PROTECTION

- **Paging** mechanisms

- determine whether an address is valid
- enforce read/write/execute permissions (granularity: regions)
- pages in kernel memory are accessible only from kernel code

- Operating systems may enforce high-level security mechanisms to hinder memory vulnerability exploitation attacks

- DEP (Data Execution Prevention)
- ASLR (Address Space Layout Randomization)
- Heap allocation randomization & other protections in Windows 10+

MEMORY ADDRESSING

- **Addressing modes** provide a way to express the addresses of data to be read from/written to the main memory
- Expressions can take immediate operands, registers, or both
- Addressing modes enable operand combinations to compute complex expressions that depend on the program state
- Usually, an instruction can have at most one memory operand

COMMON ADDRESSING MODES

Some popular ways to specify data to be read from/written to the main memory. For the last two rows, the *offset* field is optional.

Mode	Intel syntax
Immediate	<code>mov eax, [0x1000]</code>
Register	<code>mov eax, [esi]</code>
Register + offset	<code>mov eax, [esp-8]</code>
Register*width + offset	<code>mov eax, [ebx*4+0xff]</code>
Base + Reg*width + offset	<code>mov eax, [edx+ebx*4+8]</code>

BASIC X86 INSTRUCTIONS

- Programs are made of three kinds of instructions
 - **Data movement**
 - **Arithmetic & logic**
 - **Program flow control**
- Data movement instructions are straightforward
 - **mov** dest, src
Copy {register, memory content, immediate} to {register, memory location}
 - **push** src
pop dest
Used for stack manipulation (we will talk about them in a while...)

ARITHMETIC INSTRUCTIONS

- These instructions are relevant for code analysis for many reasons
 - computing offsets and function addresses
 - stack pointer modification
 - used also for data movement or to update EFLAGS

add dest, src	dest += src
sub dest, src	dest -= src
inc dest	++dest
dec dest	—dest
not dest	1's complement
neg dest	2's complement

LOGIC INSTRUCTIONS

- Common instances

- zero-ing registers: **xor** `eax, eax`
- extracting bits and checking conditions
- obfuscation

- Beware that these operations are **bit-wise**. High-level logical operations such as `&&` and `||` have a very different semantics

and <code>dest, src</code>	bitwise AND
or <code>dest, src</code>	bitwise OR
xor <code>dest, src</code>	bitwise XOR

ROTATE/SHIFT INSTRUCTIONS

Frequently encountered in bit manipulations:

shr dest, src shl dest, src	Unsigned shift (right/left)
sar dest, src sal dest, src	Signed shift
ror dest, src rol dest, src	Rotate
rcr dest, src rcl dest, src	Rotate with carry

(when src is absent, a default of 1 is assumed)

PROGRAM FLOW CONTROL


- Control flow can be implemented in three ways
 - **Unconditional branch**: EIP is replaced with some desired address
 - **Conditional branch**: EIP is replaced with some desired address depending on the value of one or more bits from EFLAGS
 - **Function calls and returns**: special type of unconditional branches
- Two primitives needed:
 - a way to specify the target of a branch (static or dynamic)
 - a way to evaluate a condition on the program state
- Code locations in assembly can be annotated with labels

UNCONDITIONAL BRANCHING

An unconditional branch can take as destination:

- an **offset**, specified as a relative offset from the current EIP value or as an absolute offset (from the base of the current code segment...)
- or an **absolute address**, provided as a register or memory operand

```
jmp 0x4  
jmp L  
sub eax, 0  
jmp eax  
L:  
add eax, 0  
jmp [eax]
```



0:	e9 00 00 00 00	jmp	5 <_main+0x5>
5:	eb 05	jmp	c <L>
7:	83 e8 00	sub	eax,0x0
a:	ff e0	jmp	eax
0000000c <L>:			
c:	83 c0 00	add	eax,0x0
f:	ff 20	jmp	DWORD PTR [eax]

Absolute addresses are straightforward, offsets may be not 😊 [7+5=c above]

CONDITIONAL BRANCHING

- Conditional branches evaluate conditions over selected bits of the EFLAGS register. Common cases:
 - **CF** (carry flag), meant for errors in unsigned arithmetics
 - **OF** (overflow flag), meant for errors in signed arithmetics
 - **SF** (sign flag), set when the result of an operation is negative
 - **ZF** (zero flag), set when the result of an operation is zero
- We can use the **cmp** instruction to compare two values
 - **cmp dest, src** computes the difference dest-src without modifying dest, and sets the flags above in a way that enables many useful comparisons...

CONDITIONAL BRANCHING

- A **jcc** instruction checks the condition specified by **cc** suffix and jumps to the given offset accordingly (no absolute addresses)
- Say we just executed **cmp D, S**:

	Jump to L if	Negated version
je L ; jz L	D==S	[D!=S] jne L ; jnz L
jg L ; jnle L	D>S	[D<=S] jng L; jle L
jge L ; jnl L	D>=S	[D<S] jnge L; jl L
ja L ; jnbe L	> with unsigned operands	jna L ; jbe L
jae L ; jnb L	>= with unsigned operands	jnae L ; jb L

OTHER COMPARISONS

- Another common comparison instruction is **test**
 - it computes the bitwise AND of its operands without altering them
 - say we want to check whether register D holds zero as value:
test D, D will set the zero flag iff. $D == 0$, then **jz** can be used
- But malware won't use only **cmp** and **test**
 - Say, you can decrement a value with **dec** or **sub** and then check if the instruction set the zero flag with **jz** 🤪 (also optimizing compilers do that!)

INTEL VS. AT&T SYNTAX

x86 assembly has two main syntax types: AT&T and Intel. Intel syntax is dominant in the Windows world. Relevant differences:

	Intel	AT&T
Parameter order	Destination, then source	Source, then destination
Parameter size	Derived from register that is used	Suffix (b, w, l)
Sigils	Autodetect	\$ for immediates, % for register names
Addresses	[expression] ; add byte/word/dword when ambiguous	offset(base, idx, width)

FUNCTIONS AND STACK FRAMES

- **Function** $\approx$ unit of code that controls register values and its portion of stack independently of other units
 - if a code unit calls another, the latter should not clobber registers in use
 - CPUs come with few general-purpose registers: one can *spill* (i.e., save) some values on the stack and fetch them later in the execution as needed
- The stack is (conceptually) divided into **frames**, one per currently active function. A frame usually includes:
 - the arguments for the function invocation
 - local function variables and other storage
 - return address for the call (i.e., where to resume execution upon function exit)

LOCAL STORAGE

- Suppose we don't have free registers left and we would like to add to EDX the difference between EBX and ECX:

```
mov TMP, eax
mov eax, ebx
sub eax, ecx
add edx, eax
mov eax, TMP
```



```
push eax
mov eax, ebx
sub eax, ecx
add edx, eax
pop eax
```

TMP ≈ temporary location for saving and restoring data currently in EAX, so that we can “borrow” EAX for a while



push eax



pop eax

PUSH AND POP INSTRUCTIONS

- **push src** moves the top of the stack up by 4 bytes, then copies the content of the operand (immediate or register) to the address pointed by the updated ESP
- **pop dest** copies the 4 bytes currently pointed by ESP to the dest register, then moves the top of the stack down by 4 bytes
- push/pop operations can also be realized via sub/add + mov:

```
push ebp  
xor ecx, ecx  
pop edx
```



```
sub esp, 4  
mov [esp], ebp  
xor ecx, ecx  
mov edx, [esp]  
add esp, 4
```

USING THE BASE POINTER

- When entering a function that makes use of local variables, register EBP can be used to reference their locations via fixed offsets

```
mov ebp, esp
sub esp, 8
mov [ebp-8], 0x1111
mov [ebp-4], 0xabadcafe
...
mov eax, [ebp-8]
```

Accessing local variables with ESP is like chasing a moving target. However, if code logic uses EBP differently, dereferencing ESP+offset is the only option left...



CALL AND RET INSTRUCTIONS


- A code piece may invoke a function using the **call** instruction
- Unlike jumps, before the control transfer takes place, it stores on the stack the address of the instruction that follows it
 - This address is called the return address for the call
 - The callee can fetch it using **ret** and resume execution in the caller: this instruction will populate EIP with the address read from the top of the stack
- Beware: *functions exists only at a logical level*
 - Code is laid out in memory without any separators
 - Adversaries may scramble their layout, making it hard for analysts to identify function boundaries!

FUNCTION CALL: AN EXAMPLE

```
push eax
add eax, 1
call myfun
sub eax, 1
```

```
myfun:
mov ebp, esp
add esp, 8
...
sub esp, 8
ret
```

```
push eax
add eax, 1
call myfun
sub eax, 1
```



CALLING CONVENTIONS

- (At least) two questions left to answer:
 - How are register values preserved across function invocations?
 - How are parameters being passed to functions?
- **Calling conventions** regulate:
 - how parameters are passed (via registers/stack) and in which order
 - which registers the callee must preserve for the caller
 - who performs ESP realignment upon return
- Malware may use custom calling conventions for its functions, but has to follow standard ones to interact with Windows APIs

WINDOWS CONVENTIONS




- Different conventions can be used in the same program
- **cdecl** (C programs)
 - arguments pushed on the stack, right-to-left order
 - caller saves EAX, ECX, and EDX when in use (i.e., free for callee, which has to preserve instead EBP, EBX, EDI, and ESI whenever it needs them)
 - caller adjusts ESP after return
 - 32-bit **return value** stored in EAX
- **stdcall** (Win32 API)
 - arguments, caller/callee-save registers and return value as in cdecl
 - callee adjusts the stack pointer upon return
 - **ret N** instruction is used, where N indicates to add N+4 bytes to ESP *after* EIP is updated

EXAMPLE: CDECL

```
mov ecx, 20
...
push ecx
push 10
push 30
call subtract
add esp, 8
pop ecx
cmp eax, ecx
jnz ErrorLabel
```

```
subtract:
push ebp
mov ebp, esp
mov eax, [ebp+8]
mov edx, [ebp+12]
sub eax, edx
pop ebp
ret
```

```
int __cdecl__
subtract(int arg1,
int arg2) {
    return arg1-arg2;
}
```






EXAMPLE: STDCALL

```
mov ecx, 20
...
push ecx
push 10
push 30
call subtract
add esp, 8
pop ecx
cmp eax, ecx
jnz ErrorLabel
```

```
subtract:
push ebp
mov ebp, esp
mov eax, [ebp+8]
mov edx, [ebp+12]
sub eax, edx
pop ebp
ret 8
```

```
int __stdcall__
subtract(int arg1,
int arg2) {
    return arg1-arg2;
}
```



CALLER-SAVE VS CALLEE-SAVE

- **Caller-save** registers are pushed to the stack **before a call**, if the caller is going to use their values after the call
 - typically EAX, ECX, EDX *(cdecl, stdcall)*
- **Callee-save** registers are pushed to the stack in the **prologue** of the callee when it needs to use them
 - typically EDI, ESI, EBX *(cdecl, stdcall)*
- Compilers often prefer callee-save registers for performance: when register pressure gets high, also caller-save registers are used. But in custom code we can find any sort of assortment...

EBP - ENTER AND LEAVE

- When used as base pointer a function prologue may set EBP with:

push ebp

mov ebp, esp

This can be done with a single instruction: **enter**

(optional immediate subtracted from ESP to make room for local variables)

- To leave a function, ESP must point to the caller's return address.

EBP may come to the rescue in a function epilogue:

mov esp, ebp

pop ebp

This can be done with a single instruction: **leave**

LEA INSTRUCTION

- LEA (Load Effective Address) evaluates an expression that defines a memory address and writes it to a destination register
 - data movement and ALU instructions dereference a given address and operate on the content of the pointed location
 - LEA simply computes an address!
- Two uses
 - address calculation (pointer arithmetic)
 - arithmetic expressions

LEA: AN EXAMPLE

- Suppose $EBX=0x1000$, and $0xABADCAFE$ is stored at $0x1004$
 - `mov eax, [ebx+4]` will copy $0xABADCAFE$ to register EAX
 - `lea eax, [ebx+4]` will copy $0x1004$ to register EAX
- Operands may also not be addresses
 - Suppose $EAX=10$, $EBX=4$ and we have to compute $C=A+2*B-1$
`lea ecx, [eax+2*ebx-1]` will do the trick with a single instruction!
 - Constraints: width for multiplication has to be 1, 2 or 4

REP INSTRUCTIONS

- Special instructions for manipulating bytes in buffers
 - **cmpsb** to compare bytes from two buffers
 - **stosb** to initialize all bytes in a buffer with same value
 - **movsb** to copy one buffer into another
 - **scasb** to search a byte in a buffer
- They are preceded by a **REP** prefix: it repeats the operation until $ECX == 0$ or an additional condition on ZF is verified
 - EDI contains the address of the first buffer (destination), ESI contains the address of the second buffer (source), and AL contains the byte to set/search
- We can implement memcmp, memset, memcpy and memchr from the C standard library using just the instructions above

HANDS-ON & BONUS PARTS

GEARING UP YOUR MACHINE

- We will use **IDA Freeware**, which is pre-installed in the VM for this course
- If you have your own VM, get IDA Freeware 8.3 or higher from the Hex Rays website (quick link: https://out7.hex-rays.com/files/idafree83_windows.exe)
- In that case, you may also need the following packages:
 - Get 64-bit MSVCRT from https://aka.ms/vs/17/release/vc_redist.x64.exe
 - (Optional) get 32-bit MSVCRT too https://aka.ms/vs/17/release/vc_redist.x86.exe

CALLS\CDECL.EXE

- Take a look at the C source used for generating this executable
- What are the addresses of the `main()`, `test()`, and `mysum()` functions?
- Function `test()`
 - What is the purpose of the `sub esp, 18h` instruction?
 - Why the compiler pushes ESI and EBX on the stack after the prologue?
 - What is the role of the `lea eax, [esi+ebx]` instruction?
 - What is the role of the instructions involving `[ebp-0c]`?
- Function `mysum()`
 - Can you tell why the compiler emitted only `pop ebp` instead of a `leave` (or an equivalent `mov esp, ebp; pop ebp` sequence) as function epilogue?

CALLS\POINTERS.EXE

- Take a look at the C source used for generating this executable
- Function main()
 - How do you take the address of a local variable?
 - How do you take the address of a global variable?
 - What can you tell by the looks of the two printed addresses?
- Function mysum()
 - How is the pointer dereferencing performed?

CALLS\STDCALL.EXE

- Take a look at the C source used for generating this executable
- What are the addresses of the `main()`, `test()`, and `mysum()` functions?
- What can you speculate about the calling conventions that `test()` and `mysum()` follow by just looking at their assembly code?
- Function `test()`
 - What is the purpose of the `sub esp, 28h` instruction?
 - What about the subsequent `sub esp, 0Ch` instruction?
 - What is the lowest value ESP can reach during the execution? (Hint: take a close look at all the components that may alter it...)
 - Can you think of better ways to optimize stack allocation to use fewer bytes?

MORE WINDOWS CONVENTIONS

Two less popular conventions

■ **fastcall**

- first two arguments (left-to-right) in ECX and EDX, others on stack (right-to-left)
- caller saves EAX, ECX, and EDX when necessary
- callee adjusts stack pointer on return (as in stdcall)

■ **thiscall** (Visual C++ for non-static class member functions)

- object pointer **this** is passed in ECX, arguments on stack (right-to-left)
- if the number of arguments is fixed, callee adjusts stack pointer on return (as in stdcall); when is variable, caller adjusts stack pointer (as in cdecl)

REP*.EXE

- We will take a look at instructions that compactly encode common control flow patterns
- Knowing popular usage patterns can speed up your code analysis
- For REP bear in mind that (where this applies) EDI holds a destination address, ESI a source address, and AL the byte of interest. ECX is used as counter.
- At a glance:
 - memcmp() to compare bytes using **rep cmpb**
 - memcpy() to copy buffers using **rep movsb**
 - memchr() to look up occurrences using **rep scasb**
 - memset() to initialize buffers using **rep stosb**

LOOP INSTRUCTION

- Loops are normally encoded using jump instructions. However, a special **loop** instruction is available (compilers avoid it for efficiency)
- ECX is used as accumulator: loop first decrements it, then if the value is not zero it jumps to the specified label

```
mov ecx, 10  
L: add eax, 1  
loop L
```

- For complex conditions, **loop** can be combined with a ZF check:
 - **loope** jumps when $ECX \neq 0 \ \&\& \ ZF = 1$
 - **loopne** jumps when $ECX \neq 0 \ \&\& \ ZF = 0$

ANTI-DISASSEMBLY

- `antidisassembly.exe` hinders the disassembly process two steps
 - it computes a displacement for the instruction pointer by first leaking its value using the **`call $+5`** pattern (we will meet it again in a few weeks...)
 - it jumps to the computed address, whereas the disassembler may erroneously believe it to be part of an instruction starting a few bytes earlier (*put in other words, the disassembler chose a “wrong” initial position to look for valid instructions*)
- Using a debugger clearly helps, but we can handle this one statically:
 - **`call $+5`** pushes on the stack the address of the subsequent **`pop eax`** instruction
 - this address is written to register EAX and is equal to 401468 (hexadecimal)
 - the target of the jump is $401468 + 0A = 401472$. Execution will continue there!
 - older versions of the IDA disassembler interpret it as part of the instruction starting at 401470 (*=> undefine such instruction and instruct IDA that some code starts instead at 401472; once the code is displayed, you can also create a function there*)

MORE ANTI-DISASSEMBLY

- `smc.exe` takes `antidisassembly.exe` to the next level
 - it leaks the instruction pointer value using the **`call $+5`** pattern and uses it to compute two addresses that get stored in EAX and EDX
 - the two take part in a loop that swaps memory contents (*bear in mind the code called `VirtualProtect` to allow code edits*) for endianness affecting 4 bytes at a time
 - when it reaches the **`nop`** instruction, something odd is about to happen!
- By using a debugger (or 📖 & 🖋️ if you prefer 🤪) we can tell that:
 - the **`pop ebp; ret`** sequence is gone
 - the two subsequent bytes (*hint: undefine the dword `1400EB02h` at address `401488`*) **`02 eb`** after the swap became a **`jmp $+2`** placed next to the **`nop`**
 - this will jump past the swapped **`ret ; pop ebp`** sequence and do the real gig!
 - the code that gets executed is the same that we extracted for `antidisassembly.exe`